

# Load\_Test\_Model

March 25, 2024

```
[1]: import torch
from torch import nn
from tqdm.auto import tqdm
from torchvision import transforms
from torchvision.utils import make_grid
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt
torch.manual_seed(0)

def show_tensor_images(image_tensor, num_images=25, size=(1, 28, 28)):
    '''
    Function for visualizing images: Given a tensor of images, number of
    ↪ images, and
    size per image, plots and prints the images in an uniform grid.
    '''
    image_tensor = (image_tensor + 1) / 2
    image_shifted = image_tensor
    image_unflat = image_shifted.detach().cpu().view(-1, *size)
    image_grid = make_grid(image_unflat[:num_images], nrow=5)
    plt.imshow(image_grid.permute(1, 2, 0).squeeze())
    plt.show()

import glob
import random
import os
from torch.utils.data import Dataset
from PIL import Image

# Inspired by https://github.com/aitorzip/PyTorch-CycleGAN/blob/master/datasets.py
↪ py
class ImageDataset(Dataset):
    def __init__(self, root, transform=None, mode='train'):
        self.transform = transform
        self.files_A = sorted(glob.glob(os.path.join(root, '%sA' % mode) + '/*.
        ↪ *'))
```

```

        self.files_B = sorted(glob.glob(os.path.join(root, '%sB' % mode) + '/*.*'))
        if len(self.files_A) > len(self.files_B):
            self.files_A, self.files_B = self.files_B, self.files_A
        self.new_perm()
        assert len(self.files_A) > 0, "Make sure you downloaded the horse2zebra_
images!"

    def new_perm(self):
        self.randperm = torch.randperm(len(self.files_B))[:len(self.files_A)]

    def __getitem__(self, index):
        item_A = self.transform(Image.open(self.files_A[index % len(self.
files_A)]))
        item_B = self.transform(Image.open(self.files_B[self.randperm[index]]))
        if item_A.shape[0] != 3:
            item_A = item_A.repeat(3, 1, 1)
        if item_B.shape[0] != 3:
            item_B = item_B.repeat(3, 1, 1)
        if index == len(self) - 1:
            self.new_perm()
        # Old versions of PyTorch didn't support normalization for
different-channel images
        return (item_A - 0.5) * 2, (item_B - 0.5) * 2

    def __len__(self):
        return min(len(self.files_A), len(self.files_B))

# In[2]:

class ResidualBlock(nn.Module):
    """
    ResidualBlock Class:
    Performs two convolutions and an instance normalization, the input is added
    to this output to form the residual block output.
    Values:
        input_channels: the number of channels to expect from a given input
    """
    def __init__(self, input_channels):
        super(ResidualBlock, self).__init__()
        self.conv1 = nn.Conv2d(input_channels, input_channels, kernel_size=3,
padding=1, padding_mode='reflect')
        self.conv2 = nn.Conv2d(input_channels, input_channels, kernel_size=3,
padding=1, padding_mode='reflect')

```

```

        self.instancenorm = nn.InstanceNorm2d(input_channels)
        self.activation = nn.ReLU()

    def forward(self, x):
        '''
        Function for completing a forward pass of ResidualBlock:
        Given an image tensor, completes a residual block and returns the
        ↪transformed tensor.
        Parameters:
            x: image tensor of shape (batch size, channels, height, width)
        '''
        original_x = x.clone()
        x = self.conv1(x)
        x = self.instancenorm(x)
        x = self.activation(x)
        x = self.conv2(x)
        x = self.instancenorm(x)
        return original_x + x

# In[3]:

class ContractingBlock(nn.Module):
    '''
    ContractingBlock Class
    Performs a convolution followed by a max pool operation and an optional
    ↪instance norm.
    Values:
        input_channels: the number of channels to expect from a given input
    '''
    def __init__(self, input_channels, use_bn=True, kernel_size=3,
    ↪activation='relu'):
        super(ContractingBlock, self).__init__()
        self.conv1 = nn.Conv2d(input_channels, input_channels * 2,
    ↪kernel_size=kernel_size, padding=1, stride=2, padding_mode='reflect')
        self.activation = nn.ReLU() if activation == 'relu' else nn.LeakyReLU(0.
    ↪2)

        if use_bn:
            self.instancenorm = nn.InstanceNorm2d(input_channels * 2)
            self.use_bn = use_bn

    def forward(self, x):
        '''
        Function for completing a forward pass of ContractingBlock:
        Given an image tensor, completes a contracting block and returns the
        ↪transformed tensor.

```

```

        Parameters:
            x: image tensor of shape (batch size, channels, height, width)
        """
        x = self.conv1(x)
        if self.use_bn:
            x = self.instancenorm(x)
        x = self.activation(x)
        return x

class ExpandingBlock(nn.Module):
    """
    ExpandingBlock Class:
    Performs a convolutional transpose operation in order to upsample,
    with an optional instance norm
    Values:
        input_channels: the number of channels to expect from a given input
    """
    def __init__(self, input_channels, use_bn=True):
        super(ExpandingBlock, self).__init__()
        self.conv1 = nn.ConvTranspose2d(input_channels, input_channels // 2,
        ↪kernel_size=3, stride=2, padding=1, output_padding=1)
        if use_bn:
            self.instancenorm = nn.InstanceNorm2d(input_channels // 2)
        self.use_bn = use_bn
        self.activation = nn.ReLU()

    def forward(self, x):
        """
        Function for completing a forward pass of ExpandingBlock:
        Given an image tensor, completes an expanding block and returns the
        ↪transformed tensor.
        Parameters:
            x: image tensor of shape (batch size, channels, height, width)
            skip_con_x: the image tensor from the contracting path (from the
        ↪opposing block of x)
            for the skip connection
        """
        x = self.conv1(x)
        if self.use_bn:
            x = self.instancenorm(x)
        x = self.activation(x)
        return x

class FeatureMapBlock(nn.Module):
    """
    FeatureMapBlock Class
    The final layer of a Generator -

```

```

    maps each the output to the desired number of output channels
    Values:
        input_channels: the number of channels to expect from a given input
        output_channels: the number of channels to expect for a given output
    '''
    def __init__(self, input_channels, output_channels):
        super(FeatureMapBlock, self).__init__()
        self.conv = nn.Conv2d(input_channels, output_channels, kernel_size=7,
padding=3, padding_mode='reflect')

    def forward(self, x):
        '''
        Function for completing a forward pass of FeatureMapBlock:
        Given an image tensor, returns it mapped to the desired number of
channels.
        Parameters:
            x: image tensor of shape (batch size, channels, height, width)
        '''
        x = self.conv(x)
        return x

# In[4]:

class Generator(nn.Module):
    '''
    Generator Class
    A series of 2 contracting blocks, 9 residual blocks, and 2 expanding blocks
to
    transform an input image into an image from the other class, with an
upfeature
    layer at the start and a downfeature layer at the end.
    Values:
        input_channels: the number of channels to expect from a given input
        output_channels: the number of channels to expect for a given output
    '''
    def __init__(self, input_channels, output_channels, hidden_channels=64):
        super(Generator, self).__init__()
        self.upfeature = FeatureMapBlock(input_channels, hidden_channels)
        self.contract1 = ContractingBlock(hidden_channels)
        self.contract2 = ContractingBlock(hidden_channels * 2)
        res_mult = 4
        self.res0 = ResidualBlock(hidden_channels * res_mult)
        self.res1 = ResidualBlock(hidden_channels * res_mult)
        self.res2 = ResidualBlock(hidden_channels * res_mult)
        self.res3 = ResidualBlock(hidden_channels * res_mult)

```

```

self.res4 = ResidualBlock(hidden_channels * res_mult)
self.res5 = ResidualBlock(hidden_channels * res_mult)
self.res6 = ResidualBlock(hidden_channels * res_mult)
self.res7 = ResidualBlock(hidden_channels * res_mult)
self.res8 = ResidualBlock(hidden_channels * res_mult)
self.expand2 = ExpandingBlock(hidden_channels * 4)
self.expand3 = ExpandingBlock(hidden_channels * 2)
self.downfeature = FeatureMapBlock(hidden_channels, output_channels)
self.tanh = torch.nn.Tanh()

def forward(self, x):
    """
    Function for completing a forward pass of Generator:
    Given an image tensor, passes it through the U-Net with residual blocks
    and returns the output.
    Parameters:
        x: image tensor of shape (batch size, channels, height, width)
    """
    x0 = self.upfeature(x)
    x1 = self.contract1(x0)
    x2 = self.contract2(x1)
    x3 = self.res0(x2)
    x4 = self.res1(x3)
    x5 = self.res2(x4)
    x6 = self.res3(x5)
    x7 = self.res4(x6)
    x8 = self.res5(x7)
    x9 = self.res6(x8)
    x10 = self.res7(x9)
    x11 = self.res8(x10)
    x12 = self.expand2(x11)
    x13 = self.expand3(x12)
    xn = self.downfeature(x13)
    return self.tanh(xn)

```

# In[5]:

```

class Discriminator(nn.Module):
    """
    Discriminator Class
    Structured like the contracting path of the U-Net, the discriminator will
    output a matrix of values classifying corresponding portions of the image
    as real or fake.
    Parameters:
        input_channels: the number of image input channels
    """

```

```

        hidden_channels: the initial number of discriminator convolutional
        ↪ filters
    """
    def __init__(self, input_channels, hidden_channels=64):
        super(Discriminator, self).__init__()
        self.upfeature = FeatureMapBlock(input_channels, hidden_channels)
        self.contract1 = ContractingBlock(hidden_channels, use_bn=False,
        ↪ kernel_size=4, activation='lrelu')
        self.contract2 = ContractingBlock(hidden_channels * 2, kernel_size=4,
        ↪ activation='lrelu')
        self.contract3 = ContractingBlock(hidden_channels * 4, kernel_size=4,
        ↪ activation='lrelu')
        self.final = nn.Conv2d(hidden_channels * 8, 1, kernel_size=1)

    def forward(self, x):
        x0 = self.upfeature(x)
        x1 = self.contract1(x0)
        x2 = self.contract2(x1)
        x3 = self.contract3(x2)
        xn = self.final(x3)
        return xn

# In[6]:

import torch.nn.functional as F

adv_criterion = nn.MSELoss()
recon_criterion = nn.L1Loss()

n_epochs = 20
dim_A = 3
dim_B = 3
display_step = 200
batch_size = 1
lr = 0.0002
load_shape = 286
target_shape = 256
device = 'cuda'

# In[7]:

transform = transforms.Compose([
    transforms.Resize(load_shape),

```

```

        transforms.RandomCrop(target_shape),
        transforms.RandomHorizontalFlip(),
        transforms.ToTensor(),
    ])

import torchvision
dataset = ImageDataset("Horse2Zebra_Dataset", transform=transform)

# In[8]:

gen_AB = Generator(dim_A, dim_B)
gen_BA = Generator(dim_B, dim_A)
gen_opt = torch.optim.Adam(list(gen_AB.parameters()) + list(gen_BA.
    ↪parameters()), lr=lr, betas=(0.5, 0.999))
disc_A = Discriminator(dim_A)
disc_A_opt = torch.optim.Adam(disc_A.parameters(), lr=lr, betas=(0.5, 0.999))
disc_B = Discriminator(dim_B)
disc_B_opt = torch.optim.Adam(disc_B.parameters(), lr=lr, betas=(0.5, 0.999))

def weights_init(m):
    if isinstance(m, nn.Conv2d) or isinstance(m, nn.ConvTranspose2d):
        torch.nn.init.normal_(m.weight, 0.0, 0.02)
    if isinstance(m, nn.BatchNorm2d):
        torch.nn.init.normal_(m.weight, 0.0, 0.02)
        torch.nn.init.constant_(m.bias, 0)

# Feel free to change pretrained to False if you're training the model from
    ↪scratch
pretrained = False
if pretrained:
    pre_dict = torch.load('cycleGAN_100000.pth')
    gen_AB.load_state_dict(pre_dict['gen_AB'])
    gen_BA.load_state_dict(pre_dict['gen_BA'])
    gen_opt.load_state_dict(pre_dict['gen_opt'])
    disc_A.load_state_dict(pre_dict['disc_A'])
    disc_A_opt.load_state_dict(pre_dict['disc_A_opt'])
    disc_B.load_state_dict(pre_dict['disc_B'])
    disc_B_opt.load_state_dict(pre_dict['disc_B_opt'])
else:
    gen_AB = gen_AB.apply(weights_init)
    gen_BA = gen_BA.apply(weights_init)
    disc_A = disc_A.apply(weights_init)
    disc_B = disc_B.apply(weights_init)

```



```

# In[9]:

# UNQ_C1 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# GRADED FUNCTION: get_disc_loss
def get_disc_loss(real_X, fake_X, disc_X, adv_criterion):
    """
    Return the loss of the discriminator given inputs.
    Parameters:
        real_X: the real images from pile X
        fake_X: the generated images of class X
        disc_X: the discriminator for class X; takes images and returns real/
        ↪ fake class X
        prediction matrices
        adv_criterion: the adversarial loss function; takes the discriminator
        predictions and the target labels and returns a adversarial
        loss (which you aim to minimize)
    """
    ##### START CODE HERE #####
    pred_fake = disc_X(fake_X)
    target = torch.zeros_like(pred_fake)
    loss1 = adv_criterion(pred_fake, target)

    pred_real = disc_X(real_X)
    target = torch.ones_like(pred_real)
    loss2 = adv_criterion(pred_real, target)
    disc_loss = (loss1 + loss2) / 2
    ##### END CODE HERE #####
    return disc_loss

# In[10]:

# UNIT TEST
test_disc_X = lambda x: x * 97
test_real_X = torch.tensor(83.)
test_fake_X = torch.tensor(89.)
test_adv_criterion = lambda x, y: x * 79 + y * 73
assert torch.abs((get_disc_loss(test_real_X, test_fake_X, test_disc_X, ↵
    ↪ test_adv_criterion)) - 659054.5000) < 1e-6
test_disc_X = lambda x: x.mean(0, keepdim=True)
test_adv_criterion = torch.nn.BCEWithLogitsLoss()
test_input = torch.ones(20, 10)
# If this runs, it's a pass - checks that the shapes are treated correctly
get_disc_loss(test_input, test_input, test_disc_X, test_adv_criterion)
print("Success!")

```

```

# In[11]:

# UNQ_C2 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# GRADED FUNCTION: get_gen_adversarial_loss
def get_gen_adversarial_loss(real_X, disc_Y, gen_XY, adv_criterion):
    """
    Return the adversarial loss of the generator given inputs
    (and the generated images for testing purposes).
    Parameters:
        real_X: the real images from pile X
        disc_Y: the discriminator for class Y; takes images and returns real/
        ↪ fake class Y
        prediction matrices
        gen_XY: the generator for class X to Y; takes images and returns the
        ↪ images
        transformed to class Y
        adv_criterion: the adversarial loss function; takes the discriminator
        predictions and the target labels and returns a adversarial
        loss (which you aim to minimize)
    """
    ##### START CODE HERE #####
    fake_Y = gen_XY(real_X)
    pred_fake = disc_Y(fake_Y)
    target = torch.ones_like(pred_fake)
    adversarial_loss = adv_criterion(pred_fake, target)

    ##### END CODE HERE #####
    return adversarial_loss, fake_Y

# In[12]:

# UNIT TEST
test_disc_Y = lambda x: x * 97
test_real_X = torch.tensor(83.)
test_gen_XY = lambda x: x * 89
test_adv_criterion = lambda x, y: x * 79 + y * 73
test_res = get_gen_adversarial_loss(test_real_X, test_disc_Y, test_gen_XY,
    ↪ test_adv_criterion)
assert torch.abs(test_res[0] - 56606652) < 1e-6
assert torch.abs(test_res[1] - 7387) < 1e-6
test_disc_Y = lambda x: x.mean(0, keepdim=True)
test_adv_criterion = torch.nn.BCEWithLogitsLoss()

```

```

test_input = torch.ones(20, 10)
# If this runs, it's a pass - checks that the shapes are treated correctly
get_gen_adversarial_loss(test_input, test_disc_Y, test_gen_XY,
    ↪test_adv_criterion)
print("Success!")

# In[13]:

# UNQ_C3 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# GRADED FUNCTION: get_identity_loss
def get_identity_loss(real_X, gen_YX, identity_criterion):
    """
    Return the identity loss of the generator given inputs
    (and the generated images for testing purposes).
    Parameters:
        real_X: the real images from pile X
        gen_YX: the generator for class Y to X; takes images and returns the
    ↪images
        transformed to class X
        identity_criterion: the identity loss function; takes the real images
    ↪from X and
        those images put through a Y->X generator and returns
    ↪the identity
        loss (which you aim to minimize)
    """
    ##### START CODE HERE #####
    identity_X = gen_YX(real_X)
    identity_loss = identity_criterion(identity_X, real_X)
    ##### END CODE HERE #####
    return identity_loss, identity_X

# In[14]:

# UNIT TEST
test_real_X = torch.tensor(83.)
test_gen_YX = lambda x: x * 89
test_identity_criterion = lambda x, y: (x + y) * 73
test_res = get_identity_loss(test_real_X, test_gen_YX, test_identity_criterion)
assert torch.abs(test_res[0] - 545310) < 1e-6
assert torch.abs(test_res[1] - 7387) < 1e-6
print("Success!")

```

```

# In[15]:

# UNQ_C4 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# GRADED FUNCTION: get_cycle_consistency_loss
def get_cycle_consistency_loss(real_X, fake_Y, gen_YX, cycle_criterion):
    """
    Return the cycle consistency loss of the generator given inputs
    (and the generated images for testing purposes).
    Parameters:
        real_X: the real images from pile X
        fake_Y: the generated images of class Y
        gen_YX: the generator for class Y to X; takes images and returns the
        ↪ images
            transformed to class X
        cycle_criterion: the cycle consistency loss function; takes the real
        ↪ images from X and
            those images put through a X->Y generator and then Y->X
        ↪ generator
            and returns the cycle consistency loss (which you aim
        ↪ to minimize)
    """
    ##### START CODE HERE #####
    cycle_X = gen_YX(fake_Y)
    cycle_loss = cycle_criterion(cycle_X, real_X)
    ##### END CODE HERE #####
    return cycle_loss, cycle_X

# In[16]:

# UNIT TEST
test_real_X = torch.tensor(83.)
test_fake_Y = torch.tensor(97.)
test_gen_YX = lambda x: x * 89
test_cycle_criterion = lambda x, y: (x + y) * 73
test_res = get_cycle_consistency_loss(test_real_X, test_fake_Y, test_gen_YX,
    ↪ test_cycle_criterion)
assert torch.abs(test_res[1] - 8633) < 1e-6
assert torch.abs(test_res[0] - 636268) < 1e-6
print("Success!")

# In[17]:

```

```

# UNQ_C5 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# GRADED FUNCTION: get_gen_loss
def get_gen_loss(real_A, real_B, gen_AB, gen_BA, disc_A, disc_B, adv_criterion,
    identity_criterion, cycle_criterion, lambda_identity=0.1, lambda_cycle=10):
    """
    Return the loss of the generator given inputs.
    Parameters:
        real_A: the real images from pile A
        real_B: the real images from pile B
        gen_AB: the generator for class A to B; takes images and returns the
    images
        transformed to class B
        gen_BA: the generator for class B to A; takes images and returns the
    images
        transformed to class A
        disc_A: the discriminator for class A; takes images and returns real/
    fake class A
        prediction matrices
        disc_B: the discriminator for class B; takes images and returns real/
    fake class B
        prediction matrices
        adv_criterion: the adversarial loss function; takes the discriminator
        predictions and the true labels and returns a adversarial
        loss (which you aim to minimize)
        identity_criterion: the reconstruction loss function used for identity
    loss
        and cycle consistency loss; takes two sets of images and returns
        their pixel differences (which you aim to minimize)
        cycle_criterion: the cycle consistency loss function; takes the real
    images from X and
        those images put through a X->Y generator and then Y->X generator
        and returns the cycle consistency loss (which you aim to minimize).
        Note that in practice, cycle_criterion == identity_criterion == L1
    loss
        lambda_identity: the weight of the identity loss
        lambda_cycle: the weight of the cycle-consistency loss
    """
    # Hint 1: Make sure you include both directions - you can think of the
    generators as collaborating
    # Hint 2: Don't forget to use the lambdas for the identity loss and cycle
    loss!
    ##### START CODE HERE #####
    # Adversarial Loss -- get_gen_adversarial_loss(real_X, disc_Y, gen_XY,
    adv_criterion)
    adv_loss_AB, fake_B = get_gen_adversarial_loss(real_A, disc_B, gen_AB,
    adv_criterion) # G : A -> B

```

```

    adv_loss_BA, fake_A = get_gen_adversarial_loss(real_B, disc_A, gen_BA,
↪adv_criterion)
    # Identity Loss -- get_identity_loss(real_X, gen_YX, identity_criterion)
    identity_loss_A, identity_A = get_identity_loss(real_A, gen_BA,
↪identity_criterion)
    identity_loss_B, identity_B = get_identity_loss(real_B, gen_AB,
↪identity_criterion)
    # Cycle-consistency Loss -- get_cycle_consistency_loss(real_X, fake_Y,
↪gen_YX, cycle_criterion)
    cycle_loss_A, cycle_A = get_cycle_consistency_loss(real_A, fake_B, gen_BA,
↪cycle_criterion)
    cycle_loss_B, cycle_B = get_cycle_consistency_loss(real_B, fake_A, gen_AB,
↪cycle_criterion)
    # Total loss
    gen_loss = adv_loss_AB + adv_loss_BA + lambda_identity * (identity_loss_A +
↪identity_loss_B) + lambda_cycle * (cycle_loss_A + cycle_loss_B)
    ##### END CODE HERE #####
    return gen_loss, fake_A, fake_B

# In[18]:

# UNIT TEST
test_real_A = torch.tensor(97)
test_real_B = torch.tensor(89)
test_gen_AB = lambda x: x * 83
test_gen_BA = lambda x: x * 79
test_disc_A = lambda x: x * 47
test_disc_B = lambda x: x * 43
test_adv_criterion = lambda x, y: x * 73 + y * 71
test_recon_criterion = lambda x, y: (x + y) * 61
test_lambda_identity = 59
test_lambda_cycle = 53
test_res = get_gen_loss(
    test_real_A,
    test_real_B,
    test_gen_AB,
    test_gen_BA,
    test_disc_A,
    test_disc_B,
    test_adv_criterion,
    test_recon_criterion,
    test_recon_criterion,
    test_lambda_identity,
    test_lambda_cycle)

```

```

assert test_res[0].item() == 4047804560
assert test_res[1].item() == 7031
assert test_res[2].item() == 8051
print("Success!")

# In[19]:

from skimage import color
import numpy as np
plt.rcParams["figure.figsize"] = (10, 10)

def train(save_model=False):
    mean_generator_loss = 0
    mean_discriminator_loss = 0
    dataloader = DataLoader(dataset, batch_size=batch_size, shuffle=True)
    cur_step = 0

    for epoch in range(n_epochs):
        # Dataloader returns the batches
        # for image, _ in tqdm(dataloader):
        for real_A, real_B in tqdm(dataloader):
            # image_width = image.shape[3]
            real_A = nn.functional.interpolate(real_A, size=target_shape)
            real_B = nn.functional.interpolate(real_B, size=target_shape)
            cur_batch_size = len(real_A)
            real_A = real_A
            real_B = real_B

            ### Update discriminator A ###
            disc_A_opt.zero_grad() # Zero out the gradient before
            ↪ backpropagation
            with torch.no_grad():
                fake_A = gen_BA(real_B)
            disc_A_loss = get_disc_loss(real_A, fake_A, disc_A, adv_criterion)
            disc_A_loss.backward(retain_graph=True) # Update gradients
            disc_A_opt.step() # Update optimizer

            ### Update discriminator B ###
            disc_B_opt.zero_grad() # Zero out the gradient before
            ↪ backpropagation
            with torch.no_grad():
                fake_B = gen_AB(real_A)
            disc_B_loss = get_disc_loss(real_B, fake_B, disc_B, adv_criterion)
            disc_B_loss.backward(retain_graph=True) # Update gradients

```

```

disc_B_opt.step() # Update optimizer

### Update generator ###
gen_opt.zero_grad()
gen_loss, fake_A, fake_B = get_gen_loss(
    real_A, real_B, gen_AB, gen_BA, disc_A, disc_B, adv_criterion,
↪recon_criterion, recon_criterion
)
gen_loss.backward() # Update gradients
gen_opt.step() # Update optimizer

# Keep track of the average discriminator loss
mean_discriminator_loss += disc_A_loss.item() / display_step
# Keep track of the average generator loss
mean_generator_loss += gen_loss.item() / display_step

### Visualization code ###
if cur_step % display_step == 0:
    print(f"Epoch {epoch}: Step {cur_step}: Generator (U-Net) loss:
↪{mean_generator_loss}, Discriminator loss: {mean_discriminator_loss}")
    show_tensor_images(torch.cat([real_A, real_B]), size=(dim_A,
↪target_shape, target_shape))
    show_tensor_images(torch.cat([fake_B, fake_A]), size=(dim_B,
↪target_shape, target_shape))
    mean_generator_loss = 0
    mean_discriminator_loss = 0
    # You can change save_model to True if you'd like to save the
↪model

    if save_model:
        torch.save({
            'gen_AB': gen_AB.state_dict(),
            'gen_BA': gen_BA.state_dict(),
            'gen_opt': gen_opt.state_dict(),
            'disc_A': disc_A.state_dict(),
            'disc_A_opt': disc_A_opt.state_dict(),
            'disc_B': disc_B.state_dict(),
            'disc_B_opt': disc_B_opt.state_dict()
        }, f"cycleGAN_{cur_step}.pth")
    cur_step += 1

```

Success!  
 Success!  
 Success!  
 Success!  
 Success!



```

[6]: import os
import sys
import torch
import torchvision.transforms as transforms
from torch.autograd import Variable
from torch.utils.data import DataLoader
from torchvision.utils import save_image
from PIL import Image
import matplotlib.pyplot as plt

def test():
    # Define the generator model
    netG_A2B = Generator(input_channels=3, output_channels=3)
    netG_B2A = Generator(input_channels=3, output_channels=3)

    # Load model weights onto CPU from the specified path
    model_path = "Horse2Zebra_Trained_Model/cycleGAN_21200.pth"
    state_dict = torch.load(model_path, map_location=torch.device('cpu'))
    # Modify the state dictionary keys if necessary
    modified_state_dict = {}
    for key, value in state_dict.items():
        # Modify keys if needed to match the expected keys in your model
        modified_state_dict[new_key] = value

    netG_A2B.load_state_dict(torch.load(model_path, map_location=torch.
↪device('cpu')))
    netG_B2A.load_state_dict(torch.load(model_path, map_location=torch.
↪device('cpu')))

    # Set models to evaluation mode
    netG_A2B.eval()
    netG_B2A.eval()

    # Define CPU tensors
    Tensor = torch.FloatTensor
    input_A = Tensor(1, 3, 256, 256)
    input_B = Tensor(1, 3, 256, 256)

    # Define image transformations
    transforms_ = [transforms.Resize(256, Image.BICUBIC),
                    transforms.ToTensor(),
                    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))]

    # Create data loader for testA and testB folders
    testA_loader = DataLoader(ImageFolderDataset("Horse2Zebra_Dataset/testA",
↪transforms_=transforms_),
                               batch_size=1, shuffle=False)

```

```

    testB_loader = DataLoader(ImageFolderDataset("Horse2Zebra_Dataset/testB",
↳transforms_=transforms_),
                               batch_size=1, shuffle=False)

    # Create output directories if they don't exist
    if not os.path.exists('output/testA_Fake'):
        os.makedirs('output/testA_Fake')
    if not os.path.exists('output/testB_Fake'):
        os.makedirs('output/testB_Fake')

    # Generate fake images for testA
    for i, batch in enumerate(testA_loader):
        real_A = Variable(input_A.copy_(batch))
        fake_B = 0.5 * (netG_A2B(real_A).data + 1.0)
        file_path = f'output/testA_Fake/{i + 1}.png'
        save_image(fake_B, file_path)
        # Plot the fake images
        plt.imshow(fake_B.squeeze().permute(1, 2, 0).numpy())
        plt.show()

    # Generate fake images for testB
    for i, batch in enumerate(testB_loader):
        real_B = Variable(input_B.copy_(batch))
        fake_A = 0.5 * (netG_B2A(real_B).data + 1.0)
        file_path = f'output/testB_Fake/{i + 1}.png'
        save_image(fake_A, file_path)
        # Plot the fake images
        plt.imshow(fake_A.squeeze().permute(1, 2, 0).numpy())
        plt.show()

    # Print completion message
    sys.stdout.write('Done\n')

# Call the test function
test()

```

```

-----
NameError                                Traceback (most recent call last)
Cell In[6], line 78
    75     sys.stdout.write('Done\n')
    77 # Call the test function
--> 78 test()

Cell In[6], line 23, in test()
    20 modified_state_dict = {}
    21 for key, value in state_dict.items():
    22 # Modify keys if needed to match the expected keys in your model

```

```

---> 23     modified_state_dict[new_key] = value
      25 netG_A2B.load_state_dict(torch.load(model_path, map_location=torch.
      ↪device('cpu'))))
      26 netG_B2A.load_state_dict(torch.load(model_path, map_location=torch.
      ↪device('cpu'))))

```

**NameError:** name 'new\_key' is not defined

```

[7]: generator = Generator(input_channels=3, output_channels=3)
     discriminator = Discriminator(input_channels=3)

```

```

# Print the architecture of the Generator
print("Generator Architecture:")
print(generator)

# Print the architecture of the Discriminator
print("\nDiscriminator Architecture:")
print(discriminator)

```

Generator Architecture:

```

Generator(
  (upfeature): FeatureMapBlock(
    (conv): Conv2d(3, 64, kernel_size=(7, 7), stride=(1, 1), padding=(3, 3),
padding_mode=reflect)
  )
  (contract1): ContractingBlock(
    (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
padding_mode=reflect)
    (activation): ReLU()
    (instancenorm): InstanceNorm2d(128, eps=1e-05, momentum=0.1, affine=False,
track_running_stats=False)
  )
  (contract2): ContractingBlock(
    (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
padding_mode=reflect)
    (activation): ReLU()
    (instancenorm): InstanceNorm2d(256, eps=1e-05, momentum=0.1, affine=False,
track_running_stats=False)
  )
  (res0): ResidualBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
padding_mode=reflect)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
padding_mode=reflect)
    (instancenorm): InstanceNorm2d(256, eps=1e-05, momentum=0.1, affine=False,
track_running_stats=False)
    (activation): ReLU()
  )
)

```

```

    )
    (res1): ResidualBlock(
      (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
padding_mode=reflect)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
padding_mode=reflect)
      (instancenorm): InstanceNorm2d(256, eps=1e-05, momentum=0.1, affine=False,
track_running_stats=False)
      (activation): ReLU()
    )
    (res2): ResidualBlock(
      (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
padding_mode=reflect)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
padding_mode=reflect)
      (instancenorm): InstanceNorm2d(256, eps=1e-05, momentum=0.1, affine=False,
track_running_stats=False)
      (activation): ReLU()
    )
    (res3): ResidualBlock(
      (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
padding_mode=reflect)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
padding_mode=reflect)
      (instancenorm): InstanceNorm2d(256, eps=1e-05, momentum=0.1, affine=False,
track_running_stats=False)
      (activation): ReLU()
    )
    (res4): ResidualBlock(
      (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
padding_mode=reflect)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
padding_mode=reflect)
      (instancenorm): InstanceNorm2d(256, eps=1e-05, momentum=0.1, affine=False,
track_running_stats=False)
      (activation): ReLU()
    )
    (res5): ResidualBlock(
      (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
padding_mode=reflect)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
padding_mode=reflect)
      (instancenorm): InstanceNorm2d(256, eps=1e-05, momentum=0.1, affine=False,
track_running_stats=False)
      (activation): ReLU()
    )
    (res6): ResidualBlock(
      (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),

```

```

padding_mode=reflect)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
padding_mode=reflect)
    (instancenorm): InstanceNorm2d(256, eps=1e-05, momentum=0.1, affine=False,
track_running_stats=False)
    (activation): ReLU()
)
(res7): ResidualBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
padding_mode=reflect)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
padding_mode=reflect)
    (instancenorm): InstanceNorm2d(256, eps=1e-05, momentum=0.1, affine=False,
track_running_stats=False)
    (activation): ReLU()
)
(res8): ResidualBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
padding_mode=reflect)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
padding_mode=reflect)
    (instancenorm): InstanceNorm2d(256, eps=1e-05, momentum=0.1, affine=False,
track_running_stats=False)
    (activation): ReLU()
)
(expand2): ExpandingBlock(
    (conv1): ConvTranspose2d(256, 128, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1), output_padding=(1, 1))
    (instancenorm): InstanceNorm2d(128, eps=1e-05, momentum=0.1, affine=False,
track_running_stats=False)
    (activation): ReLU()
)
(expand3): ExpandingBlock(
    (conv1): ConvTranspose2d(128, 64, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1), output_padding=(1, 1))
    (instancenorm): InstanceNorm2d(64, eps=1e-05, momentum=0.1, affine=False,
track_running_stats=False)
    (activation): ReLU()
)
(downfeature): FeatureMapBlock(
    (conv): Conv2d(64, 3, kernel_size=(7, 7), stride=(1, 1), padding=(3, 3),
padding_mode=reflect)
)
(tanh): Tanh()
)

```

Discriminator Architecture:  
Discriminator(

```

    (upfeature): FeatureMapBlock(
      (conv): Conv2d(3, 64, kernel_size=(7, 7), stride=(1, 1), padding=(3, 3),
padding_mode=reflect)
    )
    (contract1): ContractingBlock(
      (conv1): Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1),
padding_mode=reflect)
      (activation): LeakyReLU(negative_slope=0.2)
    )
    (contract2): ContractingBlock(
      (conv1): Conv2d(128, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1),
padding_mode=reflect)
      (activation): LeakyReLU(negative_slope=0.2)
      (instancenorm): InstanceNorm2d(256, eps=1e-05, momentum=0.1, affine=False,
track_running_stats=False)
    )
    (contract3): ContractingBlock(
      (conv1): Conv2d(256, 512, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1),
padding_mode=reflect)
      (activation): LeakyReLU(negative_slope=0.2)
      (instancenorm): InstanceNorm2d(512, eps=1e-05, momentum=0.1, affine=False,
track_running_stats=False)
    )
    (final): Conv2d(512, 1, kernel_size=(1, 1), stride=(1, 1))
  )

```

[ ]: